

torch.nn

These are the basic building blocks for graphs:

torch.nn

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weight sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GP , distribute)
- Utilities
- Quantize Functions
- Lazy Modules Initialization
 - Aliases

Buffer

A kind of Tensor that should not be considered a model parameter.

Parameter

A kind of Tensor that is to be considered a module parameter.

UninitializedParameter

A parameter that is not initialized.

UninitializedBuffer

A buffer that is not initialized.

Containers

Module

Base class for all neural network modules.

Sequential

A sequential container.

ModuleList

Holds submodules in a list.

ModuleDict

Holds submodules in a dictionary.

ParameterList

Holds parameters in a list.

ParameterDict

Holds parameters in a dictionary.

Global Hooks For Module

register_module_forward_pre_hook

Register a forward pre-hook common to all modules.

register_module_forward_hook

Register a global forward hook for all the modules.

register_module_backward_hook

Register a backward hook common to all the modules.

register_module_full_backward_pre_hook

Register a backward pre-hook common to all the modules.

register_module_full_backward_hook

Register a backward hook common to all the modules.

register_module_buffer_registration_hook

Register a buffer registration hook common to all modules.

register_module_module_registration_hook

Register a module registration hook common to all modules.

register_module_parameter_registration_hook

Register a parameter registration hook common to all modules.

Convolution Layers

<code>nn.Conv1d</code>	Applies a 1D convolution over an input signal compose of several input planes.
<code>nn.Conv2d</code>	Applies a 2D convolution over an input signal compose of several input planes.
<code>nn.Conv3d</code>	Applies a 3D convolution over an input signal compose of several input planes.
<code>nn.ConvTranspose1d</code>	Applies a 1D transpose convolution operator over an input image compose of several input planes.
<code>nn.ConvTranspose2d</code>	Applies a 2D transpose convolution operator over an input image compose of several input planes.
<code>nn.ConvTranspose3d</code>	Applies a 3D transpose convolution operator over an input image compose of several input planes.
<code>nn.LazyConv1d</code>	A <code>torch.nn.Conv1d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.LazyConv2d</code>	A <code>torch.nn.Conv2d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.LazyConv3d</code>	A <code>torch.nn.Conv3d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.LazyConvTranspose1d</code>	A <code>torch.nn.ConvTranspose1d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.LazyConvTranspose2d</code>	A <code>torch.nn.ConvTranspose2d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.LazyConvTranspose3d</code>	A <code>torch.nn.ConvTranspose3d</code> module with lazy initialization of the <code>in_channels</code> argument.
<code>nn.Unfold</code>	Extracts sliding local blocks from a batched input tensor.
<code>nn.Fold</code>	Combines an array of sliding local blocks into a large containing tensor.

Pooling layers

<code>nn.MaxPool1d</code>	Applies a 1D max pooling over an input signal compose of several input planes.
<code>nn.MaxPool2d</code>	Applies a 2D max pooling over an input signal compose of several input planes.
<code>nn.MaxPool3d</code>	Applies a 3D max pooling over an input signal compose of several input planes.
<code>nn.MaxUnpool1d</code>	Computes a partial inverse of <code>MaxPool1d</code> .
<code>nn.MaxUnpool2d</code>	Computes a partial inverse of <code>MaxPool2d</code> .
<code>nn.MaxUnpool3d</code>	Computes a partial inverse of <code>MaxPool3d</code> .
<code>nn.AvgPool1d</code>	Applies a 1D average pooling over an input signal compose of several input planes.
<code>nn.AvgPool2d</code>	Applies a 2D average pooling over an input signal compose of several input planes.
<code>nn.AvgPool3d</code>	Applies a 3D average pooling over an input signal compose of several input planes.
<code>nn.FractionalMaxPool2d</code>	Applies a 2D fractional max pooling over an input signal compose of several input planes.
<code>nn.FractionalMaxPool3d</code>	Applies a 3D fractional max pooling over an input signal compose of several input planes.
<code>nn.LPPool1d</code>	Applies a 1D power-average pooling over an input signal compose of several input planes.
<code>nn.LPPool2d</code>	Applies a 2D power-average pooling over an input signal compose of several input planes.
<code>nn.LPPool3d</code>	Applies a 3D power-average pooling over an input signal compose of several input planes.
<code>nn.AdaptiveMaxPool1d</code>	Applies a 1D adaptive max pooling over an input signal compose of several input planes.

<code>nn.AdaptiveMaxPool3d</code>	Applies a 3D adaptive max pooling over an input signal composed of several input planes.
<code>nn.AdaptiveAvgPool1d</code>	Applies a 1D adaptive average pooling over an input signal composed of several input planes.
<code>nn.AdaptiveAvgPool2d</code>	Applies a 2D adaptive average pooling over an input signal composed of several input planes.
<code>nn.AdaptiveAvgPool3d</code>	Applies a 3D adaptive average pooling over an input signal composed of several input planes.

Padding Layers

<code>nn.ReflectionPad1d</code>	Pads the input tensor using the reflection of the input boundary.
<code>nn.ReflectionPad2d</code>	Pads the input tensor using the reflection of the input boundary.
<code>nn.ReflectionPad3d</code>	Pads the input tensor using the reflection of the input boundary.
<code>nn.ReplicationPad1d</code>	Pads the input tensor using replication of the input boundary.
<code>nn.ReplicationPad2d</code>	Pads the input tensor using replication of the input boundary.
<code>nn.ReplicationPad3d</code>	Pads the input tensor using replication of the input boundary.
<code>nn.ZeroPad1d</code>	Pads the input tensor boundaries with zero.
<code>nn.ZeroPad2d</code>	Pads the input tensor boundaries with zero.
<code>nn.ZeroPad3d</code>	Pads the input tensor boundaries with zero.
<code>nn.ConstantPad1d</code>	Pads the input tensor boundaries with a constant value.
<code>nn.ConstantPad2d</code>	Pads the input tensor boundaries with a constant value.
<code>nn.ConstantPad3d</code>	Pads the input tensor boundaries with a constant value.
<code>nn.CircularPad1d</code>	Pads the input tensor using circular padding of the input boundary.
<code>nn.CircularPad2d</code>	Pads the input tensor using circular padding of the input boundary.
<code>nn.CircularPad3d</code>	Pads the input tensor using circular padding of the input boundary.

Non-linear Activations (weighted sum, nonlinearity)

<code>nn.ELU</code>	Applies the Exponential Linear Unit (ELU) function, element-wise.
<code>nn.Hardshrink</code>	Applies the Hard Shrinkage (Hard shrink) function element-wise.
<code>nn.Hardsigmoid</code>	Applies the Hard sigmoid function element-wise.
<code>nn.Hardtanh</code>	Applies the Hard Tanh function element-wise.
<code>nn.Hardswish</code>	Applies the Hard swish function, element-wise.
<code>nn.LeakyReLU</code>	Applies the LeakyReLU function element-wise.
<code>nn.LogSigmoid</code>	Applies the Logsigmoid function element-wise.
<code>nn.MultiheadAttention</code>	Allows the model to jointly attend to information from different representation subspaces.
<code>nn.PReLU</code>	Applies the element-wise PReLU function.
<code>nn.ReLU</code>	Applies the rectified linear unit function element-wise.
<code>nn.ReLU6</code>	Applies the ReLU6 function element-wise.
<code>nn.RReLU</code>	Applies the randomized leaky rectified linear unit function, element-wise.
<code>nn.SELU</code>	Applies the SELU function element-wise.
<code>nn.CELU</code>	Applies the CELU function element-wise.

<code>nn.Sigmoid</code>	Applies the Sigmoid function element-wise.
<code>nn.SiLU</code>	Applies the Sigmoid Linear Unit (SiLU) function, element-wise.
<code>nn.Mish</code>	Applies the Mish function, element-wise.
<code>nn.Softplus</code>	Applies the Softplus function element-wise.
<code>nn.Softshrink</code>	Applies the soft shrinkage function element-wise.
<code>nn.Softsign</code>	Applies the element-wise Softsign function.
<code>nn.Tanh</code>	Applies the Hyperbolic Tangent (Tanh) function element-wise.
<code>nn.Tanhshrink</code>	Applies the element-wise Tanhshrink function.
<code>nn.Threshold</code>	Thresholds each element of the input Tensor.
<code>nn.GLU</code>	Applies the gated linear unit function.

Non-linear Activations (other)

<code>nn.Softmin</code>	Applies the Softmin function to an n-dimensional input Tensor.
<code>nn.Softmax</code>	Applies the Softmax function to an n-dimensional input Tensor.
<code>nn.Softmax2d</code>	Applies SoftMax over features to each spatial location.
<code>nn.LogSoftmax</code>	Applies the $\log(\text{Softmax}(x))$ function to an n-dimensional input Tensor.
<code>nn.AdaptiveLogSoftmaxWithLoss</code>	Efficient softmax approximation.

Normalization Layers

<code>nn.BatchNorm1d</code>	Applies Batch Normalization over a 2D or 3D input.
<code>nn.BatchNorm2d</code>	Applies Batch Normalization over a 4D input.
<code>nn.BatchNorm3d</code>	Applies Batch Normalization over a 5D input.
<code>nn.LazyBatchNorm1d</code>	A <code>torch.nn.BatchNorm1d</code> module with lazy initialization.
<code>nn.LazyBatchNorm2d</code>	A <code>torch.nn.BatchNorm2d</code> module with lazy initialization.
<code>nn.LazyBatchNorm3d</code>	A <code>torch.nn.BatchNorm3d</code> module with lazy initialization.
<code>nn.GroupNorm</code>	Applies Group Normalization over a mini-batch of inputs.
<code>nn.SyncBatchNorm</code>	Applies Batch Normalization over a N-Dimensional input.
<code>nn.InstanceNorm1d</code>	Applies Instance Normalization.
<code>nn.InstanceNorm2d</code>	Applies Instance Normalization.
<code>nn.InstanceNorm3d</code>	Applies Instance Normalization.
<code>nn.LazyInstanceNorm1d</code>	A <code>torch.nn.InstanceNorm1d</code> module with lazy initialization of the <code>num_features</code> argument.
<code>nn.LazyInstanceNorm2d</code>	A <code>torch.nn.InstanceNorm2d</code> module with lazy initialization of the <code>num_features</code> argument.
<code>nn.LazyInstanceNorm3d</code>	A <code>torch.nn.InstanceNorm3d</code> module with lazy initialization of the <code>num_features</code> argument.
<code>nn.LayerNorm</code>	Applies Layer Normalization over a mini-batch of inputs.
<code>nn.LocalResponseNorm</code>	Applies local response normalization over an input signal.
<code>nn.RMSNorm</code>	Applies Root Mean Square Layer Normalization over a mini-batch of inputs.

Recurrent Layers



linearity to an input sequence.

`nn.LSTM`

Apply a multi-layer long short-term memory (LSTM) RNN to an input sequence.

`nn.GRU`

Apply a multi-layer gate recurrent unit (GRU) RNN to an input sequence.

`nn.RNNCell`

An Elman RNN cell with tanh or ReLU non-linearity.

`nn.LSTMCell`

A long short-term memory (LSTM) cell.

`nn.GRUCell`

A gate recurrent unit (GRU) cell.

Transformer Layers

`nn.Transformer`

A transformer model.

`nn.TransformerEncoder`

TransformerEncoder is a stack of N encoder layers.

`nn.TransformerDecoder`

TransformerDecoder is a stack of N decoder layers.

`nn.TransformerEncoderLayer`

TransformerEncoderLayer is made up of self-attn and feed forward network.

`nn.TransformerDecoderLayer`

TransformerDecoderLayer is made up of self-attn, multi-head-attn and feed forward network.

Linear Layers

`nn.Identity`

A placeholder identity operator that is argument-insensitive.

`nn.Linear`

Applies an affine linear transformation to the incoming data: $y = xA^T + b$.

`nn.Bilinear`

Applies a bilinear transformation to the incoming data: $y = x_1^T Ax_2 + b$.

`nn.LazyLinear`

A `torch.nn.Linear` module where *in_features* is inferred.

Dropout Layers

`nn.Dropout`

During training, randomly zeroes some of the elements of the input tensor with probability *p*.

`nn.Dropout1d`

Randomly zero out entire channels.

`nn.Dropout2d`

Randomly zero out entire channels.

`nn.Dropout3d`

Randomly zero out entire channels.

`nn.AlphaDropout`

Applies Alpha Dropout over the input.

`nn.FeatureAlphaDropout`

Randomly masks out entire channels.

Sparse Layers

`nn.Embedding`

A simple lookup table that stores embeddings of a fixed dictionary and size.

`nn.EmbeddingBag`

Compute sums or means of ‘bags’ of embeddings, without instantiating the intermediate embeddings.

Distance Functions

`nn.CosineSimilarity`

Returns cosine similarity between x_1 and x_2 , computed along *dim*.

`nn.PairwiseDistance`

Computes the pairwise distance between input vectors, or between columns of input matrices.

Loss Functions

`nn.L1Loss`

Creates a criterion that measures the mean absolute error (MAE) between each element in the input x and target y .

Creates a criterion that measures the mean square error

<code>nn.CrossEntropyLoss</code>	This criterion computes the cross entropy loss between input logits and target.
<code>nn.CTCLoss</code>	The Connectionist Temporal Classification loss.
<code>nn.NLLLoss</code>	The negative log likelihood loss.
<code>nn.PoissonNLLLoss</code>	Negative log likelihood loss with Poisson distribution of target.
<code>nn.GaussianNLLLoss</code>	Gaussian negative log likelihood loss.
<code>nn.KLDivLoss</code>	The Kullback-Leibler divergence loss.
<code>nn.BCELoss</code>	Creates a criterion that measures the Binary Cross Entropy between the target and the input probabilities:
<code>nn.BCEWithLogitsLoss</code>	This loss combines a <i>Sigmoid</i> layer and the <i>BCELoss</i> in one single class.
<code>nn.MarginRankingLoss</code>	Creates a criterion that measures the loss given inputs x_1 , x_2 , two 1D mini-batch or 0D <i>Tensors</i> , and a label 1D mini-batch or 0D <i>Tensor y</i> (containing 1 or -1).
<code>nn.HingeEmbeddingLoss</code>	Measures the loss given an input tensor x and a labels tensor y (containing 1 or -1).
<code>nn.MultiLabelMarginLoss</code>	Creates a criterion that optimizes a multi-class multi-classification hinge loss (margin-based loss) between input x (a 2D mini-batch <i>Tensor</i>) and output y (which is a 2D <i>Tensor</i> of target class indices).
<code>nn.HuberLoss</code>	Creates a criterion that uses a square term if the absolute element-wise error falls below delta and a delta-scale L1 term otherwise.
<code>nn.SmoothL1Loss</code>	Creates a criterion that uses a square term if the absolute element-wise error falls below delta and an L1 term otherwise.
<code>nn.SoftMarginLoss</code>	Creates a criterion that optimizes a two-class classification logistic loss between input tensor x and target tensor y (containing 1 or -1).
<code>nn.MultiLabelSoftMarginLoss</code>	Creates a criterion that optimizes a multi-label one-versus-all loss based on max-entropy, between input x and target y of size (N, C) .
<code>nn.CosineEmbeddingLoss</code>	Creates a criterion that measures the loss given input tensors x_1 , x_2 and a <i>Tensor</i> label y with values 1 or -1.
<code>nn.MultiMarginLoss</code>	Creates a criterion that optimizes a multi-class classification hinge loss (margin-based loss) between input x (a 2D mini-batch <i>Tensor</i>) and output y (which is a 1D tensor of target class indices, $0 \leq y \leq x.size(1) - 1$):
<code>nn.TripletMarginLoss</code>	Creates a criterion that measures the triplet loss given an input tensors x_1 , x_2 , x_3 and a margin with a value greater than 0.
<code>nn.TripletMarginWithDistanceLoss</code>	Creates a criterion that measures the triplet loss given input tensors a , p , and n (representing anchor, positive, and negative examples, respectively), and a nonnegative, real-value function (“distance function”) used to compute the relationship between the anchor and positive example (“positive distance”) and the anchor and negative example (“negative distance”).

Vision Layers

<code>nn.PixelShuffle</code>	Rearrange elements in a tensor according to an upscaling factor.
<code>nn.PixelUnshuffle</code>	Reverse the PixelShuffle operation.
<code>nn.Upsample</code>	Upsamples a given multi-channel 1D (temporal), 2D (spatial) or 3D (volumetric) data.
<code>nn.UpsamplingNearest2d</code>	Applies a 2D nearest neighbor upsampling to an input signal composed of several input channels.
<code>nn.UpsamplingBilinear2d</code>	Applies a 2D bilinear upsampling to an input signal composed of several input channels.

Shuffle Layers

<code>nn.ChannelShuffle</code>	Divides and rearranges the channels in a tensor.
--------------------------------	--

DataParallel Layers (multi-GPU, distributed)



`nn.parallel.DistributedDataParallel`

Implement `distributed` data parallelism based on

`torch.distributed` at module level.

Utilities

From the `torch.nn.utils` module:

Utility functions to clip parameter gradients.

`clip_grad_norm_`

Clip the gradient norm of an iterable of parameters.

`clip_grad_norm`

Clip the gradient norm of an iterable of parameters.

`clip_grad_value_`

Clip the gradients of an iterable of parameters at specified value.

`get_total_norm`

Compute the norm of an iterable of tensors.

`clip_grads_with_norm_`

Scale the gradients of an iterable of parameters given a pre-calculated total norm and desired max norm.

Utility functions to flatten and unflatten Module parameters to and from a single vector.

`parameters_to_vector`

Flatten an iterable of parameters into a single vector.

`vector_to_parameters`

Copy slices of a vector into an iterable of parameters.

Utility functions to fuse Modules with BatchNorm modules.

`fuse_conv_bn_eval`

Fuse a convolutional module and a BatchNorm module into a single, new convolutional module.

`fuse_conv_bn_weights`

Fuse convolutional module parameters and BatchNorm module parameters into new convolutional module parameters.

`fuse_linear_bn_eval`

Fuse a linear module and a BatchNorm module into a single, new linear module.

`fuse_linear_bn_weights`

Fuse linear module parameters and BatchNorm module parameters into new linear module parameters.

Utility functions to convert Module parameter memory formats.

`convert_conv2d_weight_memory_format`

Convert `memory_format` of `nn.Conv2d.weight` to `memory_format`.

`convert_conv3d_weight_memory_format`

Convert `memory_format` of `nn.Conv3d.weight` to `memory_format`. The conversion recursively applies to nested `nn.Module`, including `module`.

Utility functions to apply and remove weight normalization from Module parameters.

`weight_norm`

Apply weight normalization to a parameter in the given module.

`remove_weight_norm`

Remove the weight normalization reparameterization from a module.

`spectral_norm`

Apply spectral normalization to a parameter in the given module.

`remove_spectral_norm`

Remove the spectral normalization reparameterization from a module.

Utility functions for initializing Module parameters.

`skip_init`

Given a module class object and `args` / `kwargs`, instantiate the module without initializing parameters / buffers.

Utility classes and functions for pruning Module parameters.

`prune.BasePruningMethod`

A abstract base class for creation of new pruning techniques.

`prune.PruningContainer`

Container holding a sequence of pruning methods for iterative pruning.

`prune.Identity`

Utility pruning method that does not prune any units but generates the pruning parametrization with a mask of ones.

	ones with the lowest L1-norm.
<code>prune.RandomStructured</code>	Prune entire (currently unprune) channels in a tensor at ran om.
<code>prune.LnStructured</code>	Prune entire (currently unprune) channels in a tensor ase on their L <code>n</code> -norm.
<code>prune.CustomFromMask</code>	
<code>prune.identity</code>	Apply pruning reparametrization without pruning any units.
<code>prune.random_unstructured</code>	Prune tensor y removing ran om (currently unprune) units.
<code>prune.l1_unstructured</code>	Prune tensor y removing units with the lowest L1-norm.
<code>prune.random_structured</code>	Prune tensor y removing ran om channels along the specifie imension.
<code>prune.ln_structured</code>	Prune tensor y removing channels with the lowest L <code>n</code> -norm along the specifie imension.
<code>prune.global_unstructured</code>	Glo ally prunes tensors correspon ing to all parameters in <code>parameters</code> y applying the specifie <code>pruning_method</code> .
<code>prune.custom_from_mask</code>	Prune tensor correspon ing to parameter calle <code>name</code> in <code>module</code> y applying the pre-compute mask in <code>mask</code> .
<code>prune.remove</code>	Remove the pruning reparameterization from a mo ule an the pruning metho from the forwar hook.
<code>prune.is_pruned</code>	Check if a mo ule is prune y looking for pruning pre-hooks.

Parametrizations implemente using the new parametrization functionality in `torch.nn.utils.parameterize.register_parametrization()`.

<code>parametrizations.orthogonal</code>	Apply an orthogonal or unitary parametrization to a matrix or a atch of matrices.
<code>parametrizations.weight_norm</code>	Apply weight normalization to a parameter in the given mo ule.
<code>parametrizations.spectral_norm</code>	Apply spectral normalization to a parameter in the given mo ule.

tility functions to parametrize Tensors on existing Mo ules. Note that these functions can e use to parametrize a given Parameter or Buffer given a specific function that maps from an input space to the parametrize space. They are not parameterizations that woul transform an o ect into a parameter. See the [Parametrizations tutorial](#) for more information on how to implement your own parametrizations.

<code>parameterize.register_parametrization</code>	Register a parametrization to a tensor in a mo ule.
<code>parameterize.remove_parametrizations</code>	Remove the parametrizations on a tensor in a mo ule.
<code>parameterize.cached</code>	Context manager that ena les the caching system within parametrizations registere with <code>register_parametrization()</code> .
<code>parameterize.is_parametrized</code>	Determine if a mo ule has a parametrization.
<code>parameterize.ParametrizationList</code>	A sequential container that hol s an manages the original parameters or uffers of a parameterize <code>torch.nn.Module</code> .

tility functions to call a given Mo ule in a stateless manner.

<code>stateless.functional_call</code>	Perform a functional call on the mo ule y replacing the mo ule parameters an uffers with the provi e ones.
--	--

tility functions in other mo ules

<code>nn.utils.rnn.PackedSequence</code>	Hol s the ata an list of <code>batch_sizes</code> of a packe sequence.
<code>nn.utils.rnn.pack_padded_sequence</code>	Packs a Tensor containing pa e sequences of varia le length.
<code>nn.utils.rnn.pad_packed_sequence</code>	Pa a packe atch of varia le length sequences.
<code>nn.utils.rnn.pad_sequence</code>	Pa a list of varia le length Tensors with <code>padding_value</code> .
<code>nn.utils.rnn.pack_sequence</code>	Packs a list of varia le length Tensors.
<code>nn.utils.rnn.unpack_sequence</code>	npack Packe Sequence into a list of varia le length Tensors.
<code>nn.utils.rnn.unpack_padded_sequence</code>	npack padded Tensor into a list of varia le length Tensors.

`nn.Flatten`

Flattens a contiguous range of dims into a tensor.

`nn.Unflatten`

unflattens a tensor dim expanding it to a desired shape.

Quantize Functions

Quantization refers to techniques for performing computations and storing tensors at lower bitwidths than floating point precision. PyTorch supports both per tensor and per channel asymmetric linear quantization. To learn more how to use quantize functions in PyTorch, please refer to the [Quantization](#) documentation.

Lazy Modules Initialization

`nn.modules.lazy.LazyModuleMixin`

A mixin for modules that lazily initialize parameters, also known as “lazy modules”.

Aliases

The following are aliases to their counterparts in `torch.nn`:

`nn.modules.normalization.RMSNorm`

Applies Root Mean Square Layer Normalization over a mini-batch of inputs.

[← Previous](#)

[Next →](#)

Docs

Access comprehensive developer documentation for PyTorch

[View Docs →](#)

Tutorials


Get in-depth tutorials for beginners and advanced developers

[View Tutorials →](#)

Resources

Find development resources and get your questions answered

[View Resources →](#)

	PyTorch	Resources	Stay up to date	PyTorch Podcasts
	Get Started	Tutorials	Facebook	Subscribe
	Features	Docs	Twitter	Apple
	Ecosystem	Discuss	YouTube	Google
	Blog	GitHub Issues	LinkedIn	Amazon
	Contributing	Brand Guidelines		

[Terms](#) | [Privacy](#)

© Copyright The Linux Foundation. The PyTorch Foundation is a project of The Linux Foundation. For website terms of use, trademark policy and other policies applicable to The PyTorch Foundation please see www.linuxfoundation.org/policies/. The PyTorch Foundation supports the PyTorch open source project, which has been established as PyTorch Project a Series of LF Projects, LLC. For policies applicable to the PyTorch Project a Series of LF Projects, LLC, please see www.lfprojects.org/policies/.