

## torch.nn.functional

### Convolution functions

<code>conv1d</code>	Applies a 1D convolution over an input signal composed of several input planes.
<code>conv2d</code>	Applies a 2D convolution over an input image composed of several input planes.
<code>conv3d</code>	Applies a 3D convolution over an input image composed of several input planes.
<code>conv_transpose1d</code>	Applies a 1D transpose convolution operator over an input signal composed of several input planes, sometimes also called “econvolution”.
<code>conv_transpose2d</code>	Applies a 2D transpose convolution operator over an input image composed of several input planes, sometimes also called “econvolution”.
<code>conv_transpose3d</code>	Applies a 3D transpose convolution operator over an input image composed of several input planes, sometimes also called “econvolution”.
<code>unfold</code>	Extract sliding local blocks from a batched input tensor.
<code>fold</code>	Combine an array of sliding local blocks into a large containing tensor.

### Pooling functions

<code>avg_pool1d</code>	Applies a 1D average pooling over an input signal composed of several input planes.
<code>avg_pool2d</code>	Applies 2D average-pooling operation in $kH \times kW$ regions by step size $sH \times sW$ steps.
<code>avg_pool3d</code>	Applies 3D average-pooling operation in $kT \times kH \times kW$ regions by step size $sT \times sH \times sW$ steps.
<code>max_pool1d</code>	Applies a 1D max pooling over an input signal composed of several input planes.
<code>max_pool2d</code>	Applies a 2D max pooling over an input signal composed of several input planes.
<code>max_pool3d</code>	Applies a 3D max pooling over an input signal composed of several input planes.
<code>max_unpool1d</code>	Compute a partial inverse of <code>MaxPool1d</code> .
<code>max_unpool2d</code>	Compute a partial inverse of <code>MaxPool2d</code> .
<code>max_unpool3d</code>	Compute a partial inverse of <code>MaxPool3d</code> .
<code>lp_pool1d</code>	Apply a 1D power-average pooling over an input signal composed of several input planes.
<code>lp_pool2d</code>	Apply a 2D power-average pooling over an input signal composed of several input planes.
<code>lp_pool3d</code>	Apply a 3D power-average pooling over an input signal composed of several input planes.
<code>adaptive_max_pool1d</code>	Applies a 1D adaptive max pooling over an input signal composed of several input planes.
<code>adaptive_max_pool2d</code>	Applies a 2D adaptive max pooling over an input signal composed of several input planes.
<code>adaptive_max_pool3d</code>	Applies a 3D adaptive max pooling over an input signal composed of several input planes.
<code>adaptive_avg_pool1d</code>	Applies a 1D adaptive average pooling over an input signal composed of several input planes.



<code>adaptive_avg_pool3d</code>	Applies 3D adaptive average pooling over an input signal composed of several input planes.
<code>fractional_max_pool2d</code>	Applies 2D fractional max pooling over an input signal composed of several input planes.
<code>fractional_max_pool3d</code>	Applies 3D fractional max pooling over an input signal composed of several input planes.

## Attention Mechanisms

The `torch.nn.attention.bias` module contains attention biases that are designed to be used with `scale_dot_product_attention`.

<code>scaled_dot_product_attention</code>	<code>scale_dot_product_attention(query, key, value, attn_mask=None, dropout_p=0.0,</code>
---	--

## Non-linear activation functions

<code>threshold</code>	Apply a threshold to each element of the input Tensor.
<code>threshold_</code>	In-place version of <code>threshold()</code> .
<code>relu</code>	Applies the rectified linear unit function element-wise.
<code>relu_</code>	In-place version of <code>relu()</code> .
<code>hardtanh</code>	Applies the Hard Tanh function element-wise.
<code>hardtanh_</code>	In-place version of <code>hardtanh()</code> .
<code>hardswish</code>	Apply hard swish function, element-wise.
<code>relu6</code>	Applies the element-wise function $\text{ReLU6}(x) = \min(\max(0, x), 6)$ .
<code>elu</code>	Apply the Exponential Linear Unit (ELU) function element-wise.
<code>elu_</code>	In-place version of <code>elu()</code> .
<code>selu</code>	Applies element-wise, $\text{SELU}(x) = scale * (\max(0, x) + \min(0, \alpha * (\exp(x) - 1)))$ , with $\alpha = 1.6732632423543772848170429916717$ and $scale = 1.0507009873554804934193349852946$ .
<code>celu</code>	Applies element-wise, $\text{CELU}(x) = \max(0, x) + \min(0, \alpha * (\exp(x/\alpha) - 1))$ .
<code>leaky_relu</code>	Applies element-wise, $\text{LeakyReLU}(x) = \max(0, x) + negative\_slope * \min(0, x)$
<code>leaky_relu_</code>	In-place version of <code>leaky_relu()</code> .
<code>prelu</code>	Applies element-wise the function $\text{PReLU}(x) = \max(0, x) + weight * \min(0, x)$ where weight is a learnable parameter.
<code>rrelu</code>	Randomize leaky ReLU.
<code>rrelu_</code>	In-place version of <code>rrelu()</code> .
<code>glu</code>	The gated linear unit.
<code>gelu</code>	When the approximate argument is 'none', it applies element-wise the function $\text{GELU}(x) = x * \Phi(x)$
<code>logsigmoid</code>	Applies element-wise $\text{LogSigmoid}(x_i) = \log\left(\frac{1}{1+\exp(-x_i)}\right)$
<code>hardshrink</code>	Applies the hard shrinkage function element-wise
<code>tanhshrink</code>	Applies element-wise, $\text{Tanhshrink}(x) = x - \text{Tanh}(x)$
<code>softsign</code>	Applies element-wise, the function $\text{SoftSign}(x) = \frac{x}{1+ x }$

<code>softmin</code>	Apply a softmin function.
<code>softmax</code>	Apply a softmax function.
<code>softshrink</code>	Applies the soft shrinkage function elementwise
<code>gumbel_softmax</code>	Sample from the Gumbel-Softmax distribution ( <a href="#">Link 1</a> <a href="#">Link 2</a> ) and optionally discretize.
<code>log_softmax</code>	Apply a softmax followed by a logarithm.
<code>tanh</code>	Applies element-wise, $\text{Tanh}(x) = \tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$
<code>sigmoid</code>	Applies the element-wise function $\text{Sigmoid}(x) = \frac{1}{1 + \exp(-x)}$
<code>hardsigmoid</code>	Apply the Hard sigmoid function element-wise.
<code>silu</code>	Apply the Sigmoid Linear Unit (SiLU) function, element-wise.
<code>mish</code>	Apply the Mish function, element-wise.
<code>batch_norm</code>	Apply Batch Normalization for each channel across a batch of data.
<code>group_norm</code>	Apply Group Normalization for last certain number of dimensions.
<code>instance_norm</code>	Apply Instance Normalization independently for each channel in every data sample within a batch.
<code>layer_norm</code>	Apply Layer Normalization for last certain number of dimensions.
<code>local_response_norm</code>	Apply local response normalization over an input signal.
<code>rms_norm</code>	Apply Root Mean Square Layer Normalization.
<code>normalize</code>	Perform $L_p$ normalization of inputs over specific dimension.

## Linear functions

<code>linear</code>	Applies a linear transformation to the incoming data: $y = xA^T + b$ .
<code>bilinear</code>	Applies a bilinear transformation to the incoming data: $y = x_1^T Ax_2 + b$

## Dropout functions

<code>dropout</code>	During training, randomly zeroes some elements of the input tensor with probability <code>p</code> .
<code>alpha_dropout</code>	Apply alpha dropout to the input.
<code>feature_alpha_dropout</code>	Randomly masks out entire channels (a channel is a feature map).
<code>dropout1d</code>	Randomly zero out entire channels (a channel is a 1D feature map).
<code>dropout2d</code>	Randomly zero out entire channels (a channel is a 2D feature map).
<code>dropout3d</code>	Randomly zero out entire channels (a channel is a 3D feature map).

## Sparse functions

<code>embedding</code>	Generate a simple lookup table that looks up embeddings in a fixed dictionary of size.
------------------------	--

one\_hot

Takes LongTensor with index values of shape  $(*)$  and returns a tensor of shape  $(*, \text{num\_classes})$  that have zeros everywhere except where the index of last dimension matches the corresponding value of the input tensor, in which case it will be 1.

Distance functions

pairwise\_distance

See `torch.nn.PairwiseDistance` for details

cosine\_similarity

Returns cosine similarity between  $x1$  and  $x2$ , compute along dim.

pdist

Computes the p-norm distance between every pair of row vectors in the input.

Loss functions

binary\_cross\_entropy

Measure Binary Cross Entropy between the target and input probabilities.

binary\_cross\_entropy\_with\_logits

Calculate Binary Cross Entropy between target and input logits.

poisson\_nll\_loss

Poisson negative log likelihood loss.

cosine\_embedding\_loss

See `CosineEmbeddingLoss` for details.

cross\_entropy

Compute the cross entropy loss between input logits and target.

ctc\_loss

Apply the Connectionist Temporal Classification loss.

gaussian\_nll\_loss

Gaussian negative log likelihood loss.

hinge\_embedding\_loss

See `HingeEmbeddingLoss` for details.

kl\_div

Compute the KL Divergence loss.

l1\_loss

Function that takes the mean element-wise absolute value difference.

mse\_loss

Measures the element-wise mean square error, with optional weighting.

margin\_ranking\_loss

See `MarginRankingLoss` for details.

multilabel\_margin\_loss

See `MultiLabelMarginLoss` for details.

multilabel\_soft\_margin\_loss

See `MultiLabelSoftMarginLoss` for details.

multi\_margin\_loss

See `MultiMarginLoss` for details.

nll\_loss

Compute the negative log likelihood loss.

huber\_loss

Computes the Huber loss, with optional weighting.

smooth\_l1\_loss

Compute the Smooth L1 loss.

soft\_margin\_loss

See `SoftMarginLoss` for details.

triplet\_margin\_loss

Compute the triplet loss between given input tensors and a margin greater than 0.

triplet\_margin\_with\_distance\_loss

Compute the triplet margin loss for input tensors using a custom distance function.

Vision functions

pixel\_shuffle

Rearranges elements in a tensor of shape  $(*, C \times r^2, H, W)$  to a tensor of shape  $(*, C, H \times r, W \times r)$ , where r is the `upscale_factor`.

<code>pad</code>	Pad a tensor.
<code>interpolate</code>	Down/up samples the input.
<code>upsample</code>	Upsample input.
<code>upsample_nearest</code>	Upsamples the input, using nearest neighbors' pixel values.
<code>upsample_bilinear</code>	Upsamples the input, using bilinear upsampling.
<code>grid_sample</code>	Compute grid sample.
<code>affine_grid</code>	Generate 2D or 3D flow field (sampling grid), given a batch of affine matrices <code>theta</code> .

## DataParallel functions (multi-GPUs, distributed)

### data\_parallel

<code>torch.nn.parallel.data_parallel</code>	Evaluate module(input) in parallel across the GPUs given in device_ids.
--	---

[← Previous](#)

[Next →](#)

### Docs

Access comprehensive developer documentation for PyTorch

[View Docs →](#)

### Tutorials


Get in-depth tutorials for beginners and advanced developers

[View Tutorials →](#)

### Resources

Find development resources and get your questions answered

[View Resources →](#)



## PyTorch

- [Get Started](#)
- [Features](#)
- [Ecosystem](#)
- [Blog](#)
- [Contributing](#)

## Resources

- [Tutorials](#)
- [Docs](#)
- [Discuss](#)
- [GitHub Issues](#)
- [Brand Guidelines](#)

## Stay updated

- [Facebook](#)
- [Twitter](#)
- [YouTube](#)
- [LinkedIn](#)

## PyTorch Podcasts

- [Spotify](#)
- [Apple](#)
- [Google](#)
- [Amazon](#)

[Terms](#) | [Privacy](#)

© Copyright The Linux Foundation. The PyTorch Foundation is a project of The Linux Foundation. For website terms of use, trademark policy and other policies applicable to The PyTorch Foundation, please see [www.linuxfoundation.org/policies/](https://www.linuxfoundation.org/policies/). The PyTorch Foundation supports the PyTorch open source project, which has been established as PyTorch Project a Series of LF Projects, LLC. For policies applicable to the PyTorch Project a Series of LF Projects, LLC, please see [www.lfprojects.org/policies/](https://www.lfprojects.org/policies/).